What Additive Risk-Adjusted Rewards Optimise in Reinforcement Learning for Asset Allocation

Matin Tavakoli

Working draft — 22 September 2026

Abstract

We study the policy induced by additive quadratic wealth-increment rewards relative to terminal mean–variance equilibrium in a controlled discrete-time market. We establish a fixed-horizon local coefficient for the irreducible mismatch of the scalar-penalty family, construct a terminal-unit reward that exactly recovers equilibrium under independent return moments, and derive the missing continuation hedge under observed Markov dependence. Exact and finite-grid comparisons adapt rewards from financial RL papers, while a separate five-seed learning study quantifies optimisation error. We find strong exact alignment improvements but no general improvement in initial mean–variance value, constrained performance or learned financial performance. Existing mean–variance policy iteration reproduces the zero-interest calibration, so this formula is not claimed as a new algorithm. Novelty of the quantitative results remains subject to further prior-art verification.

Status. Working draft. The scoped proofs are established; publication priority, unrestricted DSR, constraints and costs, and full algorithm replications remain open. Evidence labels distinguish proved, numerically verified and preliminary results.

1 Introduction and scope

1.1 Scope and reproducibility

This draft replaces the earlier unqualified surrogate-gap narrative. The paper and proposal share saved experiment data, exact assumptions, figure captions and a version-specific literature ledger. Proofs below establish mathematical claims within the model; they do not certify publication priority. The raw experiment bundle contains all seeds, exact benchmark values and training histories.

Evaluation object	**Meaning**
Precommitted policy	Initial-date maximiser of J.
Equilibrium policy	No profitable one-step deviation by any later self.
Reward optimum	Exact optimiser of the chosen additive criterion.
Learned policy	Finite-budget output, evaluated as a deterministic affine policy.

2 What the closest literature actually does

The gap is not that financial RL ignores objectives or time inconsistency. Several important papers address them directly. The relevant question is how a specified reward's optimal policy differs from an independently specified terminal-wealth target.

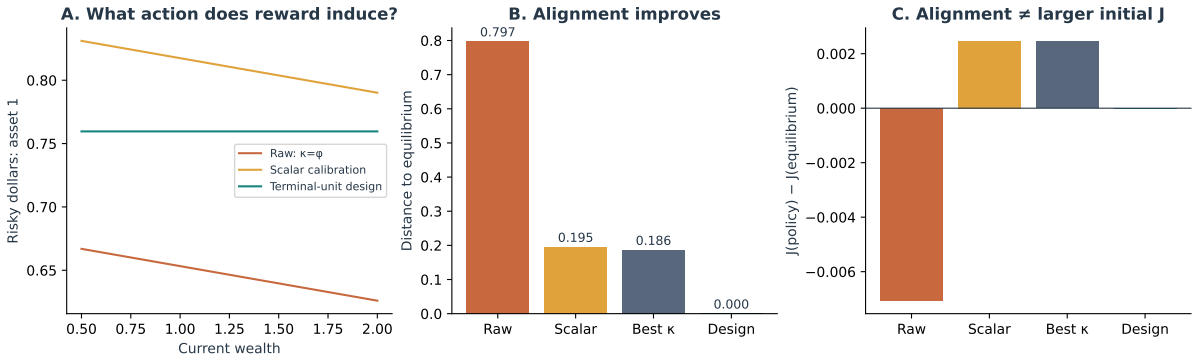


Figure 1: NUMERICALLY VERIFIED. $T = 4$, $\varphi = 1$, $s = 1.02$, three risky assets. Reward redesign eliminates exact policy mismatch. Scalar calibration attains a slightly larger initial J than equilibrium: better alignment and larger initial J are different goals.

Paper / exact locator	**Objective and method**	**Role in this project**
Kolm–Ritter (2019), manuscript Eqs. 8–10	Approximate quadratic increment reward; fitted SARSA hedging.	Closest reward family; explicitly acknowledges covariance/approximation.
Jiang et al. (2017), §4.2 Eqs. 21–22	Log-growth reward; direct gradient portfolio learning.	Alternative preference; not intrinsically a wrong growth reward.
Yang et al. (2020), §3.3 Eq. 4; Table 2	Net value change; PPO/A2C/DDPG; Sharpe ensemble selection.	Return-reward adaptation; separates reward from evaluation criterion.
Moody–Saffell (2001), §II-D Eqs. 14–15	Differential Sharpe; recurrent direct reinforcement.	Augmented-state DSR comparison, not fixed-penalty equivalence.
Bisi et al. (2020), §§3–5; Zhang et al. (2021), Algorithm 1	Per-step mean-volatility; TRVO / MVPI-TD3.	Direct prior reward-transformation theory; MVPI adaptation added.
Wang–Zhou (2020), §2 Eq. 11; Dai et al. (2023), working §4	Objective-specific MV RL; precommitment / log-MV equilibrium.	Positive precedents with analytic references; no first-MV-RL claim.

Concrete positioning: Kolm–Ritter optimise an explicitly approximate quadratic increment reward for hedging. Jiang et al. choose log growth. Yang et al. train on value increments and select agents by Sharpe. Our benchmark holds the market and feasible set fixed, solves these adapted reward problems, and compares their induced policies with a prescribed mean–variance equilibrium. This tests preference alignment; it does not refute each paper's own stated objective. (Kolm and Ritter, 2019; Jiang et al., 2017; Yang et al., 2020)

The zero-interest calibration is not a sufficient novelty claim: the finite-horizon iid reduction of Zhang's MVPI has the same fixed point. Basak–Chabakauri also discuss recovering time-consistent objectives with the same policy, and Björk–Murgoci's abstract states an associated-control result. The candidate contribution is the explicit quantitative obstruction and its audited removal within a concrete reward family. Priority remains OPEN. (Zhang et al., 2021; Basak and Chabakauri, 2010; Björk and Murgoci, 2014)

The source ledger records authors, exact titles, venue, identifiers, reward, algorithm, data, reported results, and version-specific locators for every comparison. Eleven primary versions were inspected; one additional source is explicitly abstract-only.

3 Objective, benchmark and calibration procedure

Established analytic scope: finite horizon T, initial wealth x_0, deterministic positive risk-free gross returns s_t, unconstrained risky dollar holdings u_t, and independent excess-return vectors P_t with means m_t and positive-definite covariances Σ_t. Gaussian sampling is used for training, but the exact affine moment results only require finite second moments. Long-only weights are treated separately.

$$\begin{aligned} X_{t+1} &= s_t X_t + P_t^\top u_t, \quad J(\pi) = \mathbb{E}[X_T] - \varphi \operatorname{Var}(X_T), \quad \varphi > 0, \\ M_t &= \Sigma_t + m_t m_t^\top, \qquad d_t = M_t^{-1} m_t, \quad \nu_t = m_t^\top d_t < 1, \\ H_t &= \prod_{j=t+1}^{T-1} s_j, \qquad\qquad u_t^{\mathrm{eq}} = \frac{\Sigma_t^{-1} m_t}{2\varphi H_t}. \end{aligned} \tag{1}$$

Li–Ng supply the precommitted reference; equilibrium control follows consistent planning. Our redesign targets equilibrium. If initial-date performance is the intended objective, use the precommitted benchmark and an appropriate terminal objective instead. (Li and Ng, 2000; Basak and Chabakauri, 2010; Wang and Zhou, 2020)

3.1 A reward-design rule with an explicit guarantee

$$Y_t = H_t P_t^\top u_t, \qquad r_t^{\mathrm{design}} = Y_t - \varphi(1-\nu_t) Y_t^2. \tag{2}$$

PROVED: conditional expected reward is strictly concave in u_t, wealth-independent, and maximised at u_t^{eq}. Independence and unconstrained actions make these pointwise maximisers globally optimal for the additive reward. Scaling gains into terminal units removes the interest-driven obstruction of the original wealth-increment reward. Time-varying moments require date-specific calibration.

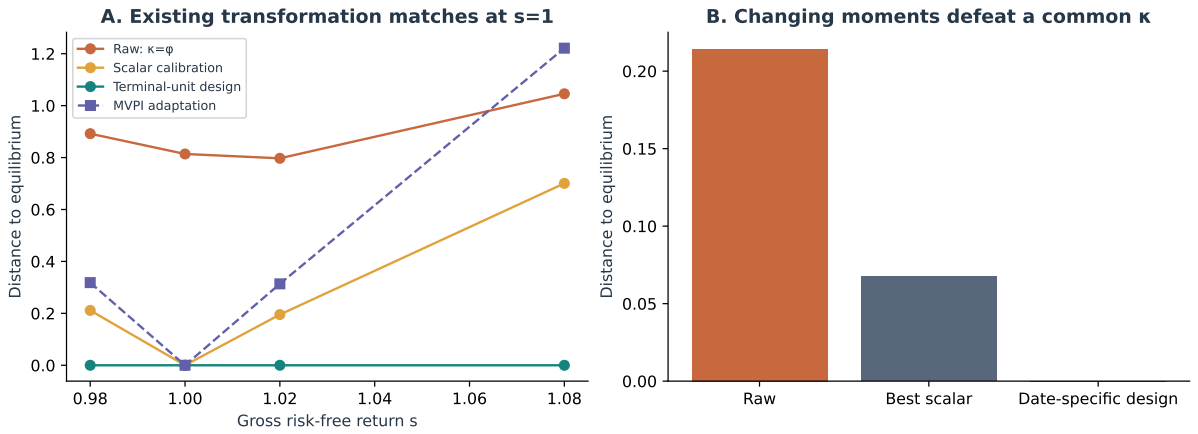


Figure 2: NUMERICALLY VERIFIED. Existing MVPI adaptation agrees at zero interest but targets per-step mean-volatility away from it. With changing moments at zero interest, a common scalar penalty leaves distance 0.067715; the date-specific design gives zero.

Implementation procedure: estimate or specify moments; declare the target and constraints; solve the exact reward problem where possible; report distance and terminal J; optimise reward parameters using a fixed reference occupancy; then train with the resulting reward and report its optimisation error separately.

4 Theoretical contribution and its boundaries

4.1 The irreducible raw-reward mismatch — PROVED

For iid returns set $d = M^{-1}m$, $\nu = m^\top d$ and $b = 1/[2\varphi(1-\nu)]$. Assume $m \neq 0$ and fixed $T \geq 2$. For $r = \Delta X - \kappa(\Delta X)^2$, let D be the square root of summed expected squared action differences along the reward policy. Then, locally around $s = 1$:

$$\begin{aligned} \min_{\kappa>0} D(\pi^\kappa, \pi^{\mathrm{eq}}) &= c_T|s-1| + O((s-1)^2), \\ c_T &= \|d\| b \sqrt{\frac{T(T^2-1)}{3} + \frac{\nu(1-\nu)T(T-1)}{2}}. \end{aligned} \tag{3}$$

At zero interest exact matching requires $\kappa = \varphi(1-\nu)$. At nonzero interest and $T \geq 2$ the final reward-optimal policy has wealth slope $-d(s-1)$, while equilibrium has zero slope and reachable wealth has positive variance. No scalar κ removes that difference. The statement concerns this reward family, not all additive rewards.

The proof closes the earlier remainder and minimisation gaps: write $h = 1/(2\kappa)$; policy intercepts and controlled wealth are affine in h, so D^2 is exactly quadratic in h. Its coefficients are analytic in $s-1$ and its unique minimum remains interior near zero. This establishes the expansion after minimisation. At large fixed-parameter T, c_T grows as $T^{3/2}$; this is not a joint uniform small-interest/large-horizon theorem.

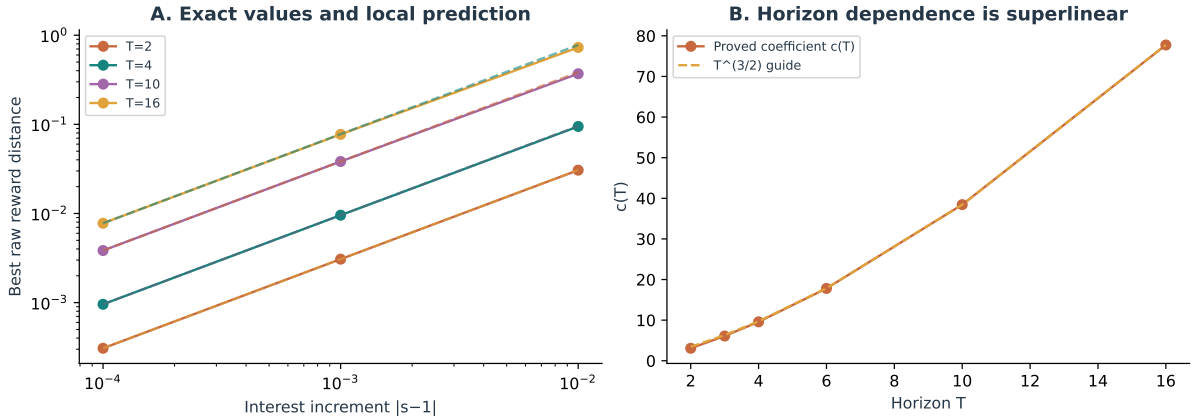


Figure 3: PROVED coefficient with NUMERICALLY VERIFIED finite-interest checks. Solid curves show exact calibration and dashed curves the local prediction. The earlier "approximately linear in horizon" claim is withdrawn.

4.2 Changing moments and a common penalty — PROVED

$$\begin{aligned} s_t = 1: \qquad \min_{\kappa>0} D^2 &= \sum_{t=0}^{T-1} a_t (b_t - \bar{b})^2, \\ a_t = \|d_t\|^2, \qquad b_t &= \frac{1}{2\varphi(1-\nu_t)}, \qquad \bar{b} = \frac{\sum_t a_t b_t}{\sum_t a_t}. \end{aligned} \tag{4}$$

The lower bound is positive exactly when b_t varies across positive-weight dates. Time variation by itself is insufficient. Degenerate zero means, the positivity condition for one-period matching, and the local nature of the theorem are retained in the proof appendix.

5 Fair comparison with published reward designs

NUMERICALLY VERIFIED: a separate long-only benchmark has $T = 3$, $x_0 = 1$, $s = 1.02$, one risky excess return $-.18$ or $.30$ with equal probability, zero costs, and a common 21-point weight grid in [0,1]. Bellman recursion exhaustively solves each finite reward tree. DSR state includes A and B. A one-step-game recursion solves the equilibrium on the same grid. The 11-point grid is also run.

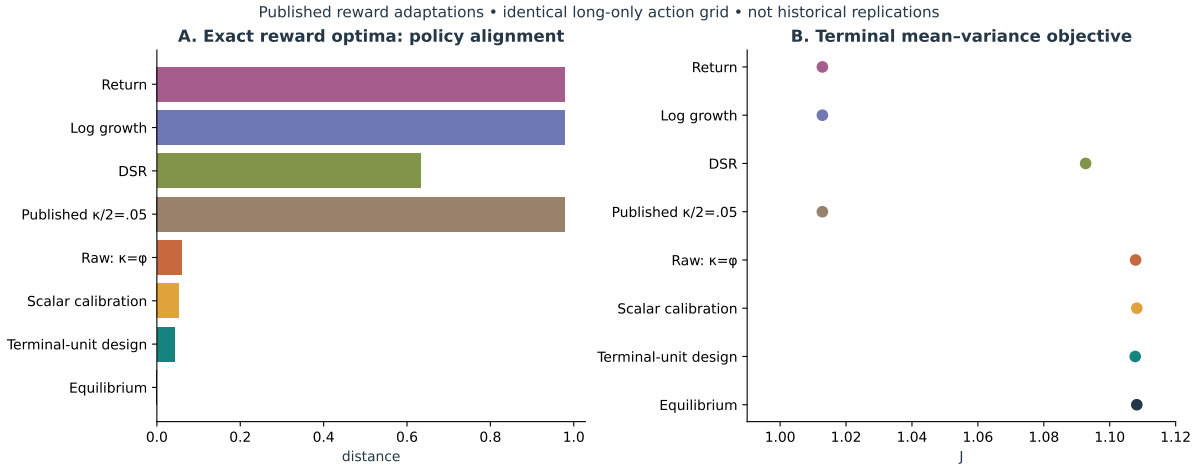


Figure 4: Adaptations of published formulas, not replications of their data, algorithms or reported performance. $\kappa/2 = .05$ is a unit-sensitive coefficient copied from the Kolm–Ritter example and is shown only as a sensitivity diagnostic.

Adapted reward	**Distance**	**Terminal J**	**Excess Sharpe**
Value change	0.97894	1.012819	0.39950
Log growth	0.97894	1.012819	0.39950
DSR	0.63257	1.092687	0.43313
Published .05 diagnostic	0.97894	1.012819	0.39950
Raw $\kappa = \varphi$	0.05848	1.107849	0.43275
Scalar calibration	0.05313	1.108217	0.43368
Terminal-unit design	0.04137	1.107729	0.43138
Grid equilibrium	0.00000	1.108204	0.43369

The design improves alignment over raw and scalar rewards in this configuration, but has slightly lower J than either. It does not exactly recover the constrained equilibrium: wealth-dependent feasible dollar holdings invalidate the unconstrained proof. Return and log-growth policies pursue different preferences, so their terminal-MV rankings are not evidence that they fail at their own tasks.

DSR uses $\eta = .1$, $A_0 = .01$, $B_0 = .01$ here. Sensitivity tests vary all three. Reward shapes are transferred; PPO/DDPG, RRL, trading costs, historical datasets and original train/test protocols are not reproduced. Algorithms cannot be ranked from this experiment.

6 Exact values, dependent returns and DSR

Independent benchmark	Policy distance	Terminal J	Excess Sharpe
Precommitted reference	3.312087	1.425799	1.171951
Equilibrium	0.000000	1.323644	0.982266
Raw $\kappa = \varphi$	0.797127	1.316580	0.987285
Scalar calibration	0.195347	1.326115	0.987292
Best scalar in raw family	0.186435	1.326093	0.987292
Terminal-unit design	0.000000	1.323644	0.982266

NUMERICALLY VERIFIED: $T = 4$, $\varphi = 1$, $s = 1.02$. Distance falls from .797127 to .195347 after zero-interest scalar calibration and to numerical zero after redesign. Raw reward has higher excess Sharpe than equilibrium (.987285 vs .982266), but lower J (1.316580 vs 1.323644). Sharpe uses excess terminal gain over $x_0 s^T$ and is not annualised.

6.1 Dependence introduces an intertemporal hedge — PROVED

$$u^{\mathrm{eq}}(t,z) = \frac{m_z}{2\varphi H_t \operatorname{Var}(P \mid z)} - \frac{\operatorname{Cov}(P, a_{t+1}(z_{t+1}) \mid z)}{H_t \operatorname{Var}(P \mid z)}. \tag{5}$$

In the observed finite-regime model, a_{t+1} is the expected terminal continuation gain. A calibration using only current conditional moments omits this covariance. Adding $-2\varphi H_t u \operatorname{Cov}(P, a_{t+1} \mid z)$ to the terminal-unit reward recovers equilibrium. This is an oracle construction using an already solved continuation, not a model-free learning result or a novel hedging principle. (Basak and Chabakauri, 2010)

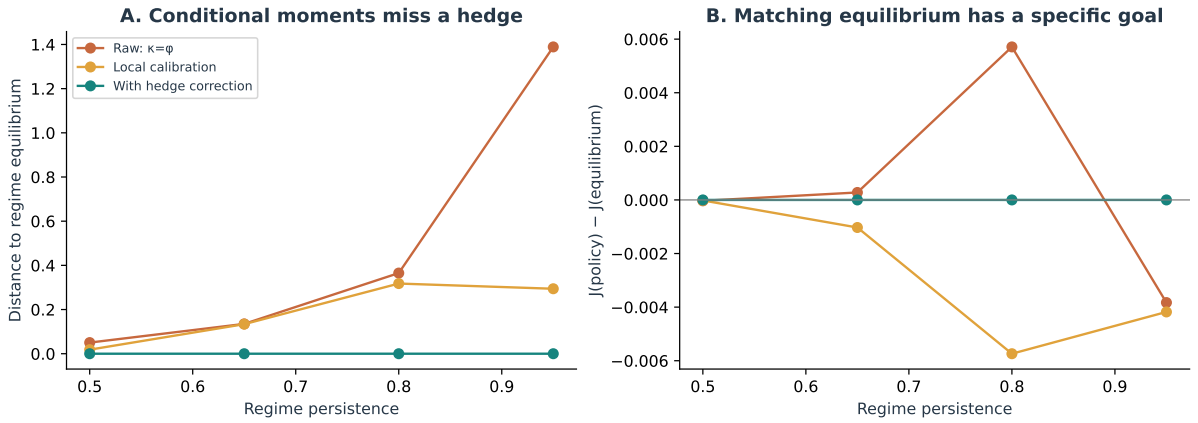


Figure 5: NUMERICALLY VERIFIED exact joint return/regime recursion. At persistence .8, raw distance is .365102, local calibration .317644, hedge-corrected reward numerical zero. Raw J exceeds equilibrium J here; that is consistent with the target being equilibrium.

R6 — PROVED: bounded finite-horizon DSR is well posed with state augmentation and $B_0 - A_0^2 > 0$. Its local quadratic identity does not establish dynamic equivalence to a fixed variance penalty. Fixed-decay EWMA has positive limiting variance; the earlier consistency claim is false. Unrestricted or infinite-horizon DSR remains OPEN.

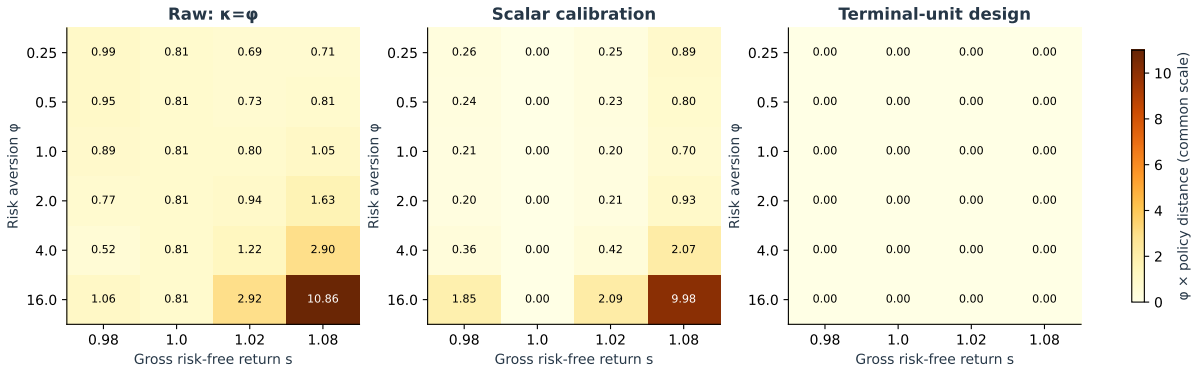


Figure 6: NUMERICALLY VERIFIED. Common-scale $\varphi \times$ distance across risk aversion and interest. The full grid has 120 parameter configurations: five horizons, six risk aversions and four rates, each with six exact reference/reward policies. Maximum redesign discrepancy is 2.35×10^{-14}.

7 Robustness and the limits of transfer

Additional exact checks vary mean, volatility and correlation over 27 markets. Redesign discrepancy remains below 5.50×10^{-15}. These verify the scoped identity, not robustness of estimated models to arbitrary market misspecification.

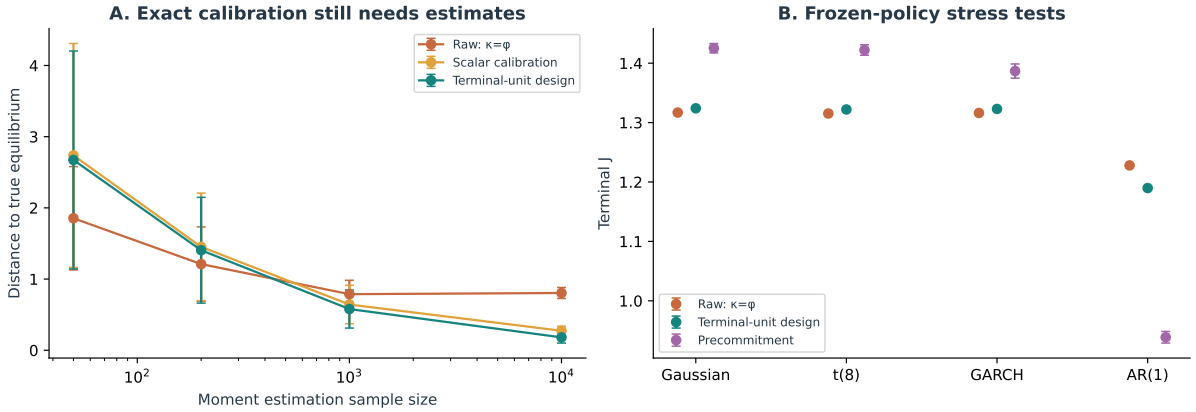


Figure 7: PRELIMINARY. Left: thirty independent moment-estimation samples at each size, mean ± sample SD. Right: frozen policies on ten batches of 50,000 paths, mean ± sample SD. GARCH and AR(1) are dependent stress tests; iid optimality is not asserted there.

The design needs accurate moments. Estimation experiments evaluate policies built from estimated moments in the true market. Stress tests hold policies fixed: matched-moment iid t(8) preserves the moment-based result, but serial correlation can reverse performance. Under AR(1), design minus raw J is −.03813 (batch SE .00027); under Gaussian iid it is +.00713 (SE .00018). No universal robustness advantage is claimed.

8 Separate learning error from reward-design error

$$
\begin{aligned}
J(\pi^{\mathrm{pre}}) - J(\widehat{\pi}) = \; & \underbrace{J(\pi^{\mathrm{pre}}) - J(\pi^{\mathrm{eq}})}_{\text{price of time consistency}} \\
& + \underbrace{J(\pi^{\mathrm{eq}}) - J(\pi^{r})}_{\text{signed reward-design difference}} \\
& + \underbrace{J(\pi^{r}) - J(\widehat{\pi})}_{\text{signed learning difference}} \quad .
\end{aligned}
\tag{6}
$$

The identity telescopes. Only the first term is guaranteed nonnegative in this benchmark. The last two are signed differences in the financial objective. Report a separate nonnegative surrogate regret: expected reward at its exact optimum minus expected reward of the learned policy. Policy distance always uses the declared reference occupancy.

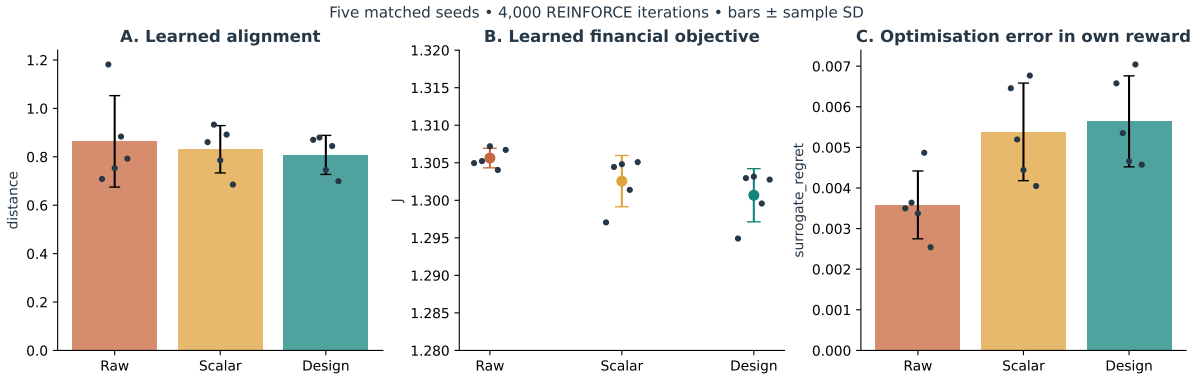


Figure 8: PRELIMINARY. Five paired seeds per reward, 4,000 iterations and batches of 512, common affine policy family, Adam learning rate .02 and exploration schedule. This is the existing REINFORCE-style implementation with normalised advantages and scaled scores. No convergence theorem is asserted.

Learned reward	**Mean distance ± SD**	**Mean J ± SD**	**Mean surrogate regret**
raw	0.8638 ± 0.1890	1.305628 ± 0.001303	0.003584
scalar	0.8312 ± 0.0978	1.302556 ± 0.003415	0.005381
design	0.8079 ± 0.0807	1.300671 ± 0.003543	0.005640

The exact redesign is aligned, but the learned redesign retains substantial distance. Its mean J is lower than the learned raw reward at this budget. Thus the current learning experiment does not support improved financial performance or broad algorithmic superiority. Reward regret is reported in each reward's own units; comparisons of its scale across rewards require care.

The fair next algorithm study must tune each reward on a separate validation budget, keep environment interactions and policy capacity comparable, add seeds, and report uncertainty and failures. The current experiment is a diagnosis, not a finished PPO/DDPG benchmark.

9 What is established, and what remains proposed

Result	Status	Boundary
R1 · Raw reward recursion and zero-interest match	PROVED	Known control tools; zero-interest result also follows from MVPI.
R2 · Fixed-horizon mismatch coefficient	PROVED	Project derivation; literature priority not established.
R3 · Terminal-unit independent-return redesign	PROVED	Unconstrained dollar actions; known moments; equilibrium target.
R4 · Common-penalty lower bound under changing moments	PROVED	Zero interest; nonzero-mean dates; weighted dispersion formula.
R5 · Observed-regime hedge and oracle correction	PROVED	Scalar risky return, exogenous observed regimes; solved continuation required.
R6 · DSR algebra and finite-horizon well-posedness	PROVED	Bounded returns, positive initial EWMA variance; no infinite-horizon theorem.
R7 · Reward adaptations and parameter sweeps	NUMERICALLY VERIFIED	Exact recursions / exact finite action trees in specified markets.
R8 · Training and estimation evidence	PRELIMINARY	Five training seeds; thirty estimation seeds; no broad superiority claim.

9.1 What can a researcher do now?

Audit an additive financial reward against a specified equilibrium, compute its exact induced policy and mismatch without learning noise, choose the best scalar penalty in the raw family, and determine whether a structural redesign is required. Under independent known moments, implement the displayed reward to recover equilibrium exactly. Under observable dependence, identify the missing continuation hedge rather than assume local calibration suffices.

9.2 Remaining research, with acceptance criteria

OPEN — Publication priority: obtain the full discrete-control prior and compare the exact coefficient and redesign with earlier inverse-control and risk-sensitive RL results. Claim no "first" result before this is done. Existing MVPI equivalence must remain in the final positioning.

OPEN — Constraints and costs: derive or numerically solve reward designs for wealth-dependent feasible sets, transaction costs and multiple dependent assets. Require a shared feasible benchmark and a quantified discretisation error before claiming exact recovery.

OPEN — DSR: analyse unrestricted policies or an appropriately regularised infinite-horizon formulation. A bounded finite-grid calculation cannot establish a constant implied terminal risk aversion or a global DSR theorem.

PROPOSED — Learning and practical validation: expand tuned policy-gradient, PPO and DDPG comparisons, implement estimated continuation hedges, and test with a preregistered historical protocol. Success means improved target alignment with controlled learning error; Sharpe and raw return remain secondary, separately reported outcomes.

The former general-horizon conjecture is resolved locally. The broad global Lambda inequality is withdrawn.

10 Evidence that changes the interpretation

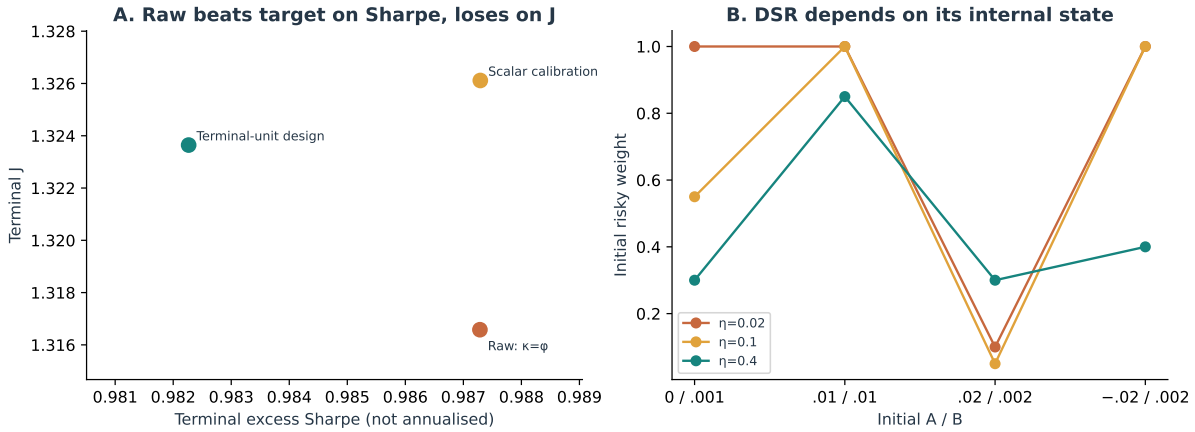


Figure 9: NUMERICALLY VERIFIED. The Sharpe/J ranking disagreement survives the corrected risk-free subtraction. DSR initial allocation varies materially with EWMA initialisation and decay; the experiment optimises full finite-horizon continuation.

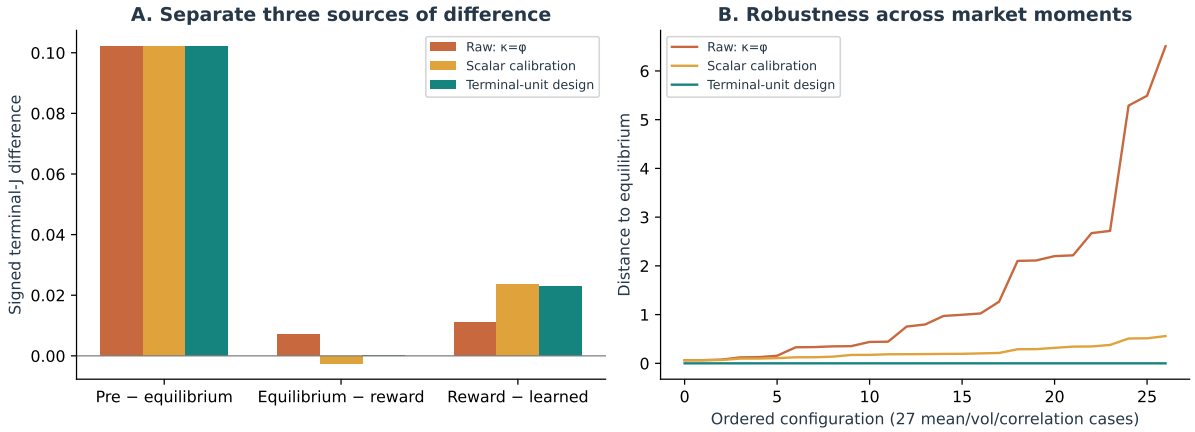


Figure 10: PRELIMINARY learning decomposition (left), NUMERICALLY VERIFIED moment sweep (right). Signed components must not be presented as three nonnegative losses. Ordered market curves compare each method's distribution of distances, not paired x-axis cases.

Main limitations: analytically simple markets and unconstrained dollar positions; moment and model knowledge; oracle continuation in the dependent redesign; small exact constrained trees; five training seeds; no transaction-cost or historical replication. These are deliberate boundaries of the evidence.

The research contribution is therefore an exact, reproducible diagnostic and a conditional design result. It is stronger than an abstract framework, but narrower than a claim that a new reward generally improves financial RL.

A Second moments and the exact reward recursion

For the wealth increment ΔX_t, the raw reward satisfies

$$\mathbb{E}[\Delta X_t - \kappa(\Delta X_t)^2] = \mathbb{E}[\Delta X_t] - \kappa \operatorname{Var}(\Delta X_t) - \kappa \mathbb{E}[\Delta X_t]^2.$$

Calling this simply a variance penalty hides an additional squared-mean penalty. The conditional-variance reward $m^\top u - \varphi u^\top \Sigma u$ has a different optimum. The zero-interest calibration follows from this distinction and the Sherman–Morrison identity; it is not an unexplained failure of Bellman optimisation.

Proposition 1 (Exact raw-reward recursion; proved)**.** *Let $M_t = \Sigma_t + m_t m_t^\top$, $d_t = M_t^{-1} m_t$ and $\nu_t = m_t^\top d_t$. For the quadratic increment reward, write the value function as $W_t(x) = p_t x^2 + q_t x + c_t$ and put $g_t = \kappa(s_t - 1) - p_{t+1} s_t$. The optimal action is $u_t(x) = \alpha_t x + \beta_t$, where*

$$\alpha_t = -\frac{d_t g_t}{\kappa - p_{t+1}}, \qquad \beta_t = \frac{d_t(1 + q_{t+1})}{2(\kappa - p_{t+1})}, \tag{7}$$

$$p_t = \frac{\kappa p_{t+1} - (1 - \nu_t) g_t^2}{\kappa - p_{t+1}}, \qquad 1 + q_t = (1 + q_{t+1}) \left(s_t - \frac{\nu_t g_t}{\kappa - p_{t+1}} \right), \tag{8}$$

with $p_T = q_T = 0$.

Proof. Completing the square in the Bellman maximisation yields the displayed formulas. The action Hessian is $-2(\kappa - p_{t+1}) M_t$. Backward induction gives $p_t \le 0$, so this Hessian is negative definite. The maximisation is therefore over all admissible Markov actions, not merely an affine search class. The constant c_t does not affect actions. □

For iid moments and $s = 1$, the optimum is $d/(2\kappa)$, whereas equilibrium is $\Sigma^{-1} m/(2\varphi) = d/[2\varphi(1 - \nu)]$. For $m \neq 0$, they agree if and only if $\kappa = \varphi(1 - \nu)$. For $m = 0$, both are zero for every $\kappa > 0$.

For $T = 1$, matching refers to the initial action, not equality of functions at all wealth values. If $m \neq 0$, a positive finite matching penalty exists precisely when

$$\frac{1}{2\varphi(1 - \nu)} + (s - 1) x_0 > 0.$$

B General-horizon small-interest mismatch

Theorem 2 (Fixed-horizon local mismatch; proved)**.** *Let $m \neq 0$ and $\Sigma \succ 0$ be constant, $T \ge 2$ fixed, $s = 1 + e$, and $b = 1/[2\varphi(1 - \nu)]$. Define the one-sided occupancy discrepancy*

$$D^2(\pi^\kappa, \pi^{\mathrm{eq}}) = \sum_{t=0}^{T-1} \mathbb{E}_{\pi^\kappa} \left\| \pi^\kappa(t, X_t) - \pi^{\mathrm{eq}}(t, X_t) \right\|^2.$$

For sufficiently small $|e|$,

$$\min_{\kappa > 0} D(\pi^\kappa, \pi^{\mathrm{eq}}) = c_T |e| + O(e^2), \tag{9}$$

$$c_T = \|d\| b \sqrt{\frac{T(T^2 - 1)}{3} + \frac{\nu(1 - \nu) T(T - 1)}{2}}. \tag{10}$$

The quantity D is a discrepancy rather than a symmetric metric. For $T = 2$, the bracket is $(2 - \nu)(1 + \nu)$. As $T \to \infty$ with other parameters fixed, $c_T \sim \|d\| b T^{3/2}/\sqrt{3}$; no uniform joint limit in e and T is asserted.

Proof. Divide p_t by κ in the recursion. The quantities p_t/κ and q_t, hence α_t, are independent of κ, while β_t is linear in $h = 1/(2\kappa)$. On each common return path, controlled wealth has the form $X_t = L_t(e)x_0 + hZ_t(e)$. Consequently $F(e,h) = D^2$ is exactly quadratic in h. Its coefficients are analytic near $e = 0$: the finite backward recursions have nonzero denominators there, and moment propagation uses only finite second moments. At zero interest,

$$F(0,h) = T\|d\|^2(h-b)^2.$$

The coefficient of h^2 is strictly positive. The global unconstrained minimiser $h^*(e)$ is therefore analytic and stays positive near zero, so it is also the global minimiser over $h > 0$. This justifies the minimisation and local Taylor expansion without assuming that a fitted local optimum is global.

Write $\kappa = \varphi(1-\nu)(1+we)$. Uniformly for bounded w,

$$\begin{aligned}
\alpha_t &= -de + O(e^2),\\
\beta_t &= db\{1 + [(T-t-1)(1-\nu) - w]e\} + O(e^2),\\
\beta_t^{\mathrm{eq}} &= db[1 - (T-t-1)e] + O(e^2).
\end{aligned}$$

At $e = 0$, the common policy is db, so

$$\mathbb{E}[X_t] = x_0 + t\nu b, \qquad \mathrm{Var}(X_t) = tb^2\nu(1-\nu).$$

The leading action difference is de times the scalar random variable

$$-X_t + b[(T-t-1)(2-\nu) - w].$$

Its means form an arithmetic progression with step $-2b$. Minimisation gives $w^* = (T-1)(1-\nu) - x_0/b$, and the centred means are $b(T-1-2t)$. Their squared sum is $b^2T(T^2-1)/3$; the sum of the variances is $b^2\nu(1-\nu)T(T-1)/2$. Analyticity yields $\min_h F(e,h) = c_T^2e^2 + O(e^3)$. Since $c_T > 0$, taking square roots gives remainder $O(e^2)$. □

The same leading coefficient holds under fixed equilibrium occupancy: at $e = 0$ the occupancies coincide and only zeroth-order moments enter the leading quadratic term. Away from zero interest, these discrepancies differ and must be labelled.

For $m \neq 0$, $s \neq 1$ and $T \geq 2$, exact matching along the surrogate occupancy is impossible. Matching at time zero forces the nonzero equilibrium action, hence positive next-wealth variance. Subsequent nonzero equilibrium actions preserve positive variance. At the final date, the surrogate slope is $-d(s-1)$ and the equilibrium slope is zero, contradicting almost-sure matching. This obstruction concerns the raw increment family, not all additive rewards.

No global bound of the previously proposed form

$$D \leq C\big(|\kappa - \kappa_0| + c|s-1|\big)$$

with C constant over all $\kappa > 0$ follows. At $s = 1$, $D = \sqrt{T}\|d\||1/(2\kappa) - b|$ diverges as $\kappa \downarrow 0$ while the proposed right-hand side remains bounded. A local bound on compact positive penalty sets is valid by smoothness.

C Time-varying independent returns and reward redesign

Let $H_t = \prod_{j=t+1}^{T-1} s_j$, with empty product one.

Proposition 3 (Independent-return exact recovery; proved)**.** *Under independent returns, deterministic positive rates, known first two moments and unconstrained risky dollar holdings, the equilibrium action is*

$$u_t^{\mathrm{eq}} = \frac{\Sigma_t^{-1} m_t}{2\varphi H_t}.$$

The additive reward

$$Y_t = H_t P_t^\top u_t, \qquad r_t^{\mathrm{design}} = Y_t - \varphi(1-\nu_t) Y_t^2$$

has this same optimal action at each date.

Proof. Under future deterministic holdings, $X_T = H_t(s_t x + P_t^\top u) +$ future gains. Independence removes action-dependent covariance with future gains. The one-step objective has action-dependent part

$$H_t m_t^\top u - \varphi H_t^2 u^\top \Sigma_t u.$$

Strict concavity gives the stated action; backward induction establishes deterministic equilibrium holdings.

Each conditional expected redesigned reward is strictly concave in u and independent of wealth. Exogenous independent returns and unconstrained actions make pointwise maximisation at every date globally optimal for the additive expectation, even among history-dependent policies. Its maximiser is

$$\frac{d_t}{2\varphi(1-\nu_t)H_t} = u_t^{\mathrm{eq}}.$$

□

The redesign measures excess gains in terminal units and applies date-specific calibration. It reproduces equilibrium, not the precommitted optimum. The alternative conditional-variance reward $Y_t - \varphi H_t^2 u_t^\top \Sigma_t u_t$ has the same conditional maximiser. Both constructions require moments and a specified target. Neither implies a universal improvement in initial J: another policy can outperform equilibrium at time zero.

If every $s_t = 1$, a common raw κ matches equilibrium if and only if all nonzero-mean dates have the same ν_t and $\kappa = \varphi(1-\nu_t)$. Writing $a_t = \|d_t\|^2$ and $b_t = 1/[2\varphi(1-\nu_t)]$, all holdings are deterministic and

$$D^2 = \sum_t a_t (h - b_t)^2, \qquad \min_h D^2 = \sum_t a_t (b_t - \bar{b})^2, \qquad \bar{b} = \frac{\sum_t a_t b_t}{\sum_t a_t}.$$

This lower bound is positive exactly when b_t varies over positive-weight dates. Time variation alone does not imply unavoidable mismatch: ν_t may remain constant. If all $a_t = 0$, every action is zero and the distance is zero; $\bar{b}$ need not be defined.

D Predictable Markov returns and the continuation hedge

For scalar excess return P and observed finite regime z, specify the joint conditional law of (P, z_{t+1}). Rates are deterministic. Let equilibrium terminal-mean continuation be $f_{t+1}(x, z') = H_t x + a_{t+1}(z')$; its conditional variance is independent of wealth. The current action-dependent variance is

$$H_t^2 \operatorname{Var}(P \mid z) u^2 + 2 H_t \operatorname{Cov}(P, a_{t+1}(z_{t+1}) \mid z) u.$$

Therefore

$$u^{\mathrm{eq}}(t,z) = \frac{m_z}{2\varphi H_t \operatorname{Var}(P \mid z)} - \frac{\operatorname{Cov}(P, a_{t+1}(z_{t+1}) \mid z)}{H_t \operatorname{Var}(P \mid z)}. \tag{11}$$

Backward propagation of conditional first and second moments establishes the formula. Positive conditional return variance makes this a strict global one-step maximum. Substituting $\kappa = \varphi(1-\nu_z)$ state by state produces only the first term. A nonzero hedge covariance defeats this substitution even at zero interest.

An equilibrium-informed reward is

$$r = Y - \varphi(1-\nu_z)Y^2 - 2\varphi \operatorname{Cov}(P, a_{t+1} \mid z) H_t u, \tag{12}$$

$$Y = H_t P u. \tag{13}$$

Its conditional derivative gives exactly u^{eq}. Exogenous transitions and wealth independence make these pointwise maximisers globally optimal for the sum. This is inverse control using an already solved continuation, not a model-free shortcut or a novelty claim. The relevant diagnostic is the quantified failure of local calibration and recovery of the missing hedge. Mean–variance hedging under dependence is established literature.

E Differential Sharpe: finite-horizon conclusions

With $A' = A + \eta(\rho - A)$, $B' = B + \eta(\rho^2 - B)$ and $V = B - A^2 > 0$,

$$V' = (1-\eta)V + \eta(1-\eta)(\rho - A)^2 > 0, \qquad 0 < \eta < 1. \tag{14}$$

For bounded returns $|\rho| \le R$, finite T, and $V_0 > 0$, $V_t \ge (1-\eta)^t V_0 > 0$, and the DSR reward is bounded on the reachable finite-horizon state set. A state including market predictors, previous holdings if costs apply, wealth, A and B supports ordinary finite-horizon Bellman recursion. No infinite-horizon fixed-point conclusion follows without additional assumptions.

The unscaled differential Sharpe signal is

$$D_{\mathrm{DSR}} = \frac{B(\rho - A) - \frac{1}{2}A(\rho^2 - B)}{V^{3/2}} = \frac{B}{V^{3/2}}\left(\rho - \frac{A}{2B}\rho^2 - \frac{A}{2}\right).$$

The legacy implementation multiplies this by η, which leaves the policy optimum unchanged for a fixed η but changes reported reward units. The displayed quadratic identity is only local: A and B depend on previous actions, the positive scale depends on state, and continuation depends on their updates. It does not imply dynamic equivalence to a fixed-κ quadratic reward. Negative A also gives a negative effective penalty. A constant implied terminal risk aversion has not been proved.

For iid ρ with finite variance σ^2, a fixed-decay EWMA obeys

$$\lim_{t\to\infty} \operatorname{Var}(A_t) = \frac{\eta\sigma^2}{2-\eta} > 0 \quad \text{when } \sigma^2 > 0.$$

Its expectation approaches the mean, but the estimator is not consistent under nondegenerate iid innovations. The earlier consistency claim is withdrawn. The corrected bounded-tree experiment optimises the augmented reward exactly on a finite action set and checks sensitivity to decay and initialisation; it does not establish an unrestricted optimum.

F Known benchmarks and the existing MVPI transformation

For iid ν, deterministic risk-free growth $S = \prod_t s_t$, and initial wealth x_0,

$$J(\pi^{\mathrm{pre}}) = Sx_0 + \frac{(1-\nu)^{-T} - 1}{4\varphi}, \qquad J(\pi^{\mathrm{eq}}) = Sx_0 + \frac{T\nu}{4\varphi(1-\nu)}. \tag{15}$$

The first follows from the Li–Ng embedding or the efficient-frontier square; the second follows by summing independent equilibrium gains in terminal units. Their difference has $1/\varphi$ scaling and is independent of x_0. These are consequences of known solutions, not new results.

Signed telescoping differences in J must be accompanied by nonnegative surrogate regret $J_r(\pi^r) - J_r(\widehat{\pi})$. Reported Sharpe is the mean terminal excess gain over $x_0 \prod_t s_t$, divided by the terminal wealth standard deviation; it is not an annualised backtest Sharpe ratio.

Zhang, Liu and Whiteson (2021), Algorithm 1, uses

$$\widehat{r} = (1 + 2\varphi y)r - \varphi r^2,$$

where y is the occupancy mean reward. In the finite-horizon uniform-time adaptation at $s = 1$ with iid returns, each exact inner solve gives deterministic identical holdings $u_n = dh_n$. Hence $y_n = m^\top u_n = \nu h_n$, and

$$h_{n+1} = \frac{1}{2\varphi} + \nu h_n \;\longrightarrow\; h^* = \frac{1}{2\varphi(1-\nu)}, \qquad 0 \le \nu < 1.$$

The fixed point is exactly the scalar-calibrated equilibrium. This is an elementary reduction of existing MVPI, not a new learning method. At nonzero interest, the implemented iteration is a fixed-point method for a per-step mean-volatility objective. Numerical convergence alone establishes neither a global optimum nor terminal mean–variance equilibrium.

References

Petter N. Kolm; Gordon Ritter. Dynamic Replication and Hedging: A Reinforcement Learning Approach. The Journal of Financial Data Science 1(1), 159–171 (2019). https://doi.org/10.3905/jfds.2019.1.1.159

Source verification. 16-page author manuscript at smallake.kr; figures are missing in this manuscript. Text and equations read; journal PDF unavailable. "Automatic hedging in theory", manuscript pp. 7–8, Eqs. (8)–(10), footnote 4; "Automatic hedging in practice", pp. 9–14, Exhibits 1–5.

Zhengyao Jiang; Dixing Xu; Jinjun Liang. A Deep Reinforcement Learning Framework for the Financial Portfolio Management Problem. arXiv preprint 1706.10059 (2017). https://arxiv.org/abs/1706.10059

Source verification. arXiv v2, 16 July 2017; saved PDF. §§2.2–2.4, §4.2 Eqs. (21)–(22), §4.3; Tables 1–2; §6.

Hongyang Yang; Xiao-Yang Liu; Shan Zhong; Anwar Walid. Deep Reinforcement Learning for Automated Stock Trading: An Ensemble Strategy. Proceedings of the First ACM International Conference on AI in Finance (ICAIF), 8 pages (2020). https://doi.org/10.1145/3383455.3422540

Source verification. Columbia-hosted author manuscript; saved PDF. Do not use the later 2025 arXiv upload as a 2020 identifier. §3.3 Eq. (4); §5 ensemble selection; §6.1 and Fig. 4; §6.2.4 and Table 2.

John Moody; Matthew Saffell. Learning to Trade via Direct Reinforcement. IEEE Transactions on Neural Networks 12(4), 875–889 (2001). https://doi.org/10.1109/72.935097

Source verification. IDSIA-hosted 17-page manuscript, visually read because text encoding is corrupted. §II-D, Eqs. (12)–(17), PDF pp. 3–4; §III-A Eq. (31); §IV-A, Fig. 5; §IV-C.2, Fig. 8.

Lorenzo Bisi; Luca Sabbioni; Edoardo Vittori; Matteo Papini; Marcello Restelli. Risk-Averse Trust Region Optimization for Reward-Volatility Reduction. Proceedings of IJCAI-20, 4583–4589 (2020). https://doi.org/10.24963/ijcai.2020/632

Source verification. Published proceedings PDF, saved and read. §3 Eqs. (3)–(6); §§4–5, Algorithm 1; §7.1–7.2, Fig. 1(a)–(d).

Shangtong Zhang; Bo Liu; Shimon Whiteson. Mean-Variance Policy Iteration for Risk-Averse Reinforcement Learning. Proceedings of AAAI 35(12), 10905–10913 (2021). https://doi.org/10.1609/aaai.v35i12.17302

Source verification. arXiv:2004.10888v6, 7 April 2022; differs in date from published 2021 proceedings. "Mean-Variance RL", Eq. (2); "Mean-Variance Policy Iteration", Eqs. (3)–(4), Algorithm 1, Propositions 1–2; "Experiments", Table 1.

Haoran Wang; Xun Yu Zhou. Continuous-Time Mean-Variance Portfolio Selection: A Reinforcement Learning Framework. Mathematical Finance 30(4), 1273–1308 (2020). https://doi.org/10.1111/mafi.12281

Source verification. arXiv:1904.11392, saved 39-page preprint; section references refer to this version. §2 Eq. (11); §3 Theorems 1 and 3; §4; §5.1 Table 1, §§5.2–5.3.

Min Dai; Yuchao Dong; Yanwei Jia. Learning Equilibrium Mean-Variance Strategy. Mathematical Finance 33(4), 1166–1212 (2023). https://doi.org/10.1111/mafi.12402

Source verification. NUS RMI Working Paper 2021-02, March 2021, 52 pages, read via web PDF. Local download blocked; journal section numbering not asserted. Working version §§2–3, §4.2 Theorem 2; §5 Algorithm 1; §6.1–6.2, Fig. 1.

Duan Li; Wan-Lung Ng. Optimal Dynamic Portfolio Selection: Multiperiod Mean-Variance Formulation. Mathematical Finance 10(3), 387–406 (2000). https://doi.org/10.1111/1467-9965.00100

Source verification. Journal PDF mirrored at tianjiablog.wordpress.com; saved. §§3–4, Theorems 1–2 (embedding); §5 riskless-asset case.

Suleyman Basak; Georgy Chabakauri. Dynamic Mean-Variance Asset Allocation. The Review of Financial Studies 23(8), 2970–3016 (2010). https://doi.org/10.1093/rfs/hhq028

Source verification. Author thesis chapter 2, "Dynamic Mean-Variance Asset Allocation", read; journal full text unavailable. Chapter numbering is not journal numbering. Thesis chapter 2 §2.2.2, Proposition 2.1, Eqs. (148)–(149); §2.2.2 Remark 1; §2.2.4; discrete-time extension in §2.4.

Andrew Y. Ng; Daishi Harada; Stuart Russell. Policy Invariance under Reward Transformations: Theory and Application to Reward Shaping. Proceedings of ICML, 278–287 (1999). https://ai.stanford.edu/~ang/papers/shaping-icml99.pdf

Source verification. Author-hosted 10-page conference paper; visually read due to encoding. No DOI asserted. §3 Theorem 1, Eq. (2), Corollary 2; §4 Figures 1–2; Appendix A.

Tomas Björk; Agatha Murgoci. A Theory of Markovian Time-Inconsistent Stochastic Control in Discrete Time. Finance and Stochastics 18(3), 545–592 (2014). https://doi.org/10.1007/s00780-014-0234-y

Source verification. Publisher metadata and abstract only; full-text verification OPEN. Not counted as a paper fully read. Publisher abstract only; no theorem number or detailed comparison asserted.