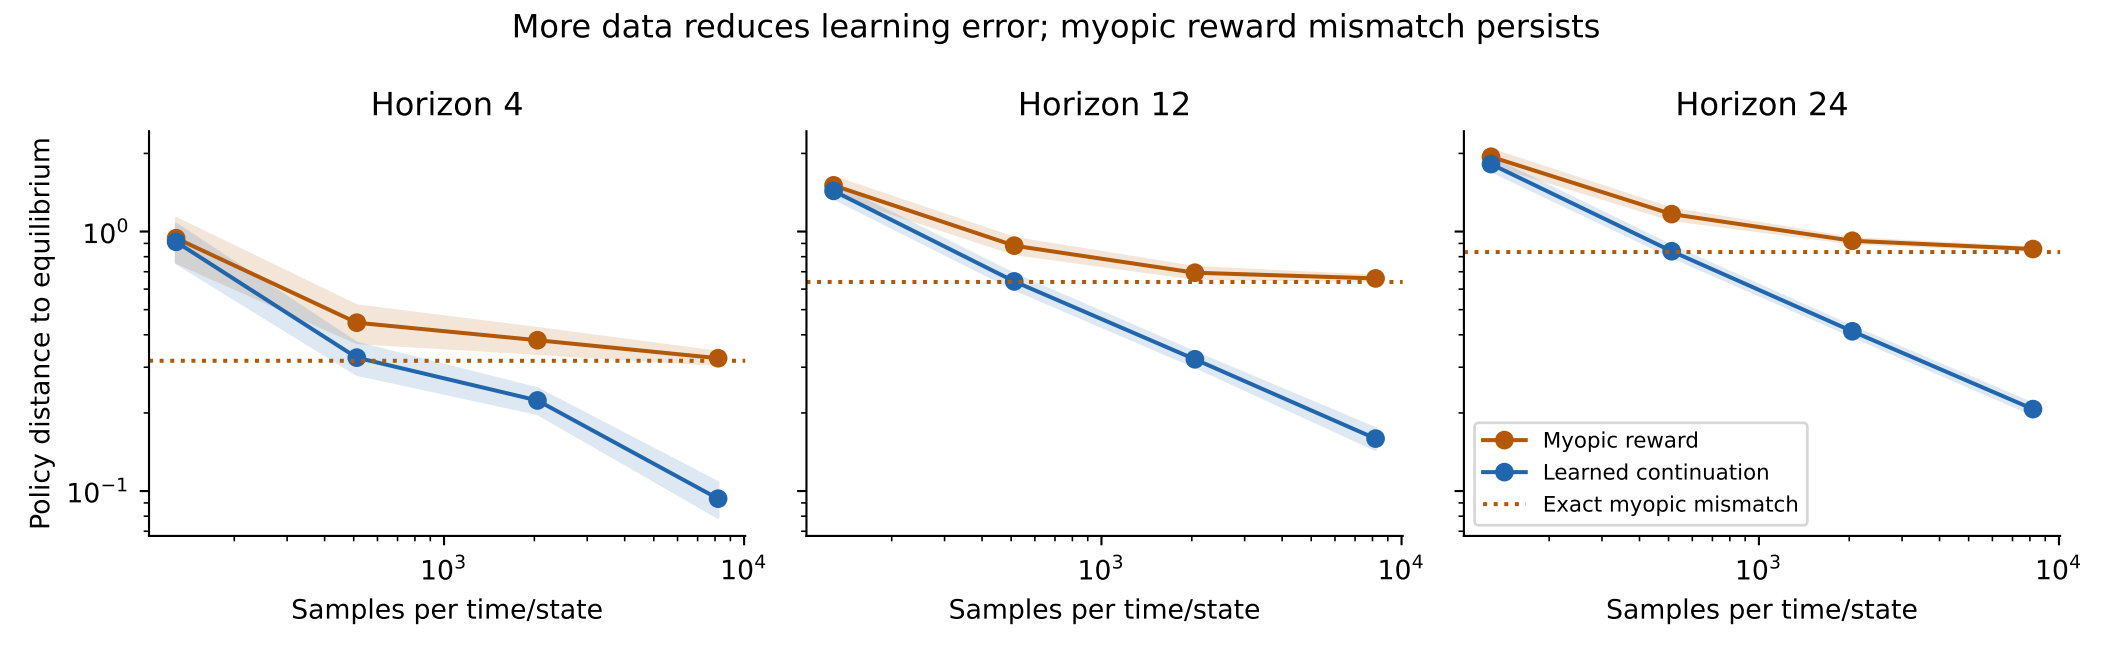
More data reduces learning error; myopic reward mismatch persists
Horizon 4
Horizon 12
Horizon 24
Policy distance to equilibrium
10^0
10^{-1}
10^3
10^4
Samples per time/state
Myopic reward
Learned continuation
Exact myopic mismatch